

Bios 6301: Assignment 5

Yudong Cao

Grade 51/50

Note: Great job. Also, it's worth learning how Cole solved Question 2 using laaply and tapply. In the future, for packages that might not be installed by collaborators, you can use the following to check for and install a package:

```
if("lubridate" %in% rownames(installed.packages()) == FALSE) {  
  install.packages("lubridate",repos="http://cran.rstudio.com/")  
}
```

Due Tuesday, 15 November, 1:00 PM

$5^{n=\text{day}}$ points taken off for each day late.

50 points total.

Submit a single knitr file (named `homework5.rmd`), along with a valid PDF output file. Inside the file, clearly indicate which parts of your responses go with which problems (you may use the original homework document as a template). Add your name as `author` to the file's metadata section. Raw R code/output or word processor files are not acceptable.

Failure to name file `homework5.rmd` or include author name may result in 5 points taken off.

Question 1

24 points

Import the HAART dataset (`haart.csv`) from the GitHub repository into R, and perform the following manipulations: (4 points each)

```
haart<-read.csv('haart.csv')
```

1. Convert date columns into a usable (for analysis) format. Use the `table` command to display the counts of the year from `init.date`.

```
haart[, 'last.visit'] <- as.POSIXct(haart[, 'last.visit'], format="%m/%d/%y")  
haart[, 'init.date'] <- as.POSIXct(haart[, 'init.date'], format="%m/%d/%y")  
haart[, 'date.death'] <- as.POSIXct(haart[, 'date.death'], format="%m/%d/%y")  
init.date.year<-sub("([0-9]{4})-([0-9]{2})-([0-9]{2})", "\\1", haart[, 'init.date'])  
table(init.date.year)
```

```
## init.date.year  
## 1998 2000 2001 2002 2003 2004 2005 2006 2007  
##    1    5   17   60  270  292  207  104   44
```

2. Create an indicator variable (one which takes the values 0 or 1 only) to represent death within 1 year of the initial visit. How many observations died in year 1?

```
time.diff<-haart[, 'date.death']-haart[, 'init.date']  
haart[, 'death.within.1year']<-ifelse(time.diff<=365,1,0)  
table(haart[, 'death.within.1year'])
```

```
##  
## 0 1  
## 25 92
```

92 observations died in year 1.

3. Use the `init.date`, `last.visit` and `death.date` columns to calculate a followup time (in days), which is the difference between the first and either the last visit or a death event (whichever comes first). If these times are longer than 1 year, censor them (this means if the value is above 365, set followup to 365). Print the quantile for this new variable.

```
fut.visit<-difftime(haart[, 'last.visit'], haart[, 'init.date'], units="days")
fut.visit<-ifelse(fut.visit>365,365,fut.visit)
fut.death<-difftime(haart[, 'date.death'], haart[, 'init.date'], units="days")
fut.death<-ifelse(fut.death>365,365,fut.death)
attach(haart)
followup.time<-0
for (i in 1:nrow(haart)) {
  if (is.na(fut.death[i])) {
    followup.time[i]=fut.visit[i]
  } else {
    followup.time[i]=min(fut.visit[i],fut.death[i])
  }
}
haart[, 'fut']<-ceiling(followup.time)
quantile(followup.time,na.rm=TRUE)
```

```
##    0%  25%  50%  75% 100%
##     0  338  365  365  365
```

JC Grading -2 Close, but this code: `followup.time[i]=min(fut.visit[i],fut.death[i])` should be `followup.time[i]=min(fut.visit[i],fut.death[i], na.rm=TRUE)`. Look at these cross-tabs:

```
table(is.na(fut.death),is.na(fut.visit))
```

```
##
##          FALSE TRUE
## FALSE    106   11
## TRUE     883    0
```

```
table(is.na(followup.time))
```

```
##
## FALSE  TRUE
##   989    11
```

4. Create another indicator variable representing loss to followup; this means the observation is not known to be dead but does not have any followup visits after the first year. How many records are lost-to-followup?

```
haart[, 'loss.tfu']<-ifelse(haart[, 'fut']==365 & haart[, 'death']==0,1,0)
table(haart[, 'loss.tfu'])
```

```
##
##    0    1
## 290 710
```

710 records are lost-to-followup.

JC Grading -2

710 overcounts the answer of 173. Since follow-up time is truncated at 365, it includes many who had visits beyond the first year. You could check if anyone had a 365 visit and if not consider ltf those whose visit was less than 365.

- Recall our work in class, which separated the `init.reg` field into a set of indicator variables, one for each unique drug. Create these fields and append them to the database as new columns. Which drug regimen are found over 100 times?

```
haart[, 'init.reg'] <- as.character(haart[, 'init.reg'])
all.reg <- strsplit(haart[, 'init.reg'], ",")
all.reg <- unlist(all.reg)
table(all.reg)
```

```
## all.reg
## 3TC ABC ATV AZT D4T DDC DDI EFV FPV FTC IDV LPV NFV NVP RTV SQV T20 TDF
## 973 56 2 794 146 1 38 516 2 8 27 31 8 358 79 29 1 10
```

3TC, AZT, D4T, EFV and NVP are found over 100 times.

```
all.reg <- unique(all.reg)
row.reg <- strsplit(haart[, 'init.reg'], ",")
user.reg <- sapply(all.reg, function(j) sapply(row.reg, function(i) j %in% i))
haart <- cbind(haart, +user.reg)
head(haart)
```

```
##   male age aids cd4baseline logvl  weight hemoglobin  init.reg
## 1    1  25   0          NA    NA      NA          NA 3TC,AZT,EFV
## 2    1  49   0         143    NA 58.0608         11 3TC,AZT,EFV
## 3    1  42   1         102    NA 48.0816          1 3TC,AZT,EFV
## 4    0  33   0         107    NA 46.0000         NA 3TC,AZT,NVP
## 5    1  27   0          52     4     NA         NA 3TC,D4T,EFV
## 6    0  34   0         157    NA 54.8856         NA 3TC,AZT,NVP
##   init.date last.visit death date.death death.within.1year fut loss.tfu
## 1 2003-07-01 2007-02-26    0      <NA>                NA 365          1
## 2 2004-11-23 2008-02-22    0      <NA>                NA 365          1
## 3 2003-04-30 2005-11-21    1 2006-01-11                0 365          0
## 4 2006-03-25 2006-05-05    1 2006-05-07                1 41          0
## 5 2004-09-01 2007-11-13    0      <NA>                NA 365          1
## 6 2003-12-02 2008-02-28    0      <NA>                NA 365          1
##   3TC AZT EFV NVP D4T ABC DDI IDV LPV RTV SQV FTC TDF DDC NFV T20 ATV FPV
## 1    1    1    1    0    0    0    0    0    0    0    0    0    0    0    0    0    0
## 2    1    1    1    0    0    0    0    0    0    0    0    0    0    0    0    0    0
## 3    1    1    1    0    0    0    0    0    0    0    0    0    0    0    0    0    0
## 4    1    1    0    1    0    0    0    0    0    0    0    0    0    0    0    0    0
## 5    1    0    1    0    1    0    0    0    0    0    0    0    0    0    0    0    0
## 6    1    1    0    1    0    0    0    0    0    0    0    0    0    0    0    0    0
```

- The dataset `haart2.csv` contains a few additional observations for the same study. Import these and append them to your master dataset (if you were smart about how you coded the previous steps, cleaning the additional observations should be easy!). Show the first five records and the last five records of the complete (and clean) data set.

```
haart2 <- read.csv('haart2.csv')
haart2[, 'last.visit'] <- as.POSIXct(haart2[, 'last.visit'], format="%m/%d/%y")
haart2[, 'init.date'] <- as.POSIXct(haart2[, 'init.date'], format="%m/%d/%y")
haart2[, 'date.death'] <- as.POSIXct(haart2[, 'date.death'], format="%m/%d/%y")
time.diff2 <- haart2[, 'date.death'] - haart2[, 'init.date']
haart2[, 'death.within.1year'] <- ifelse(time.diff2 <= 365, 1, 0)
fut.visit2 <- difftime(haart2[, 'last.visit'], haart2[, 'init.date'], units="days")
fut.visit2 <- ifelse(fut.visit2 > 365, 365, fut.visit2)
attach(haart2)
```

```
## The following objects are masked from haart:
##
##   age, aids, cd4baseline, date.death, death, death.within.1year,
##   hemoglobin, init.date, init.reg, last.visit, logvl, male,
##   weight
followup.time2<-0
for (i in 1:nrow(haart2)) {
  followup.time2[i]=fut.visit2[i]
}
haart2[, 'fut']<-ceiling(followup.time2)
haart2[, 'loss.tfu']<-ifelse(haart2[, 'fut']==365 & haart2[, 'death']==0,1,0)
haart2[, 'init.reg']<-as.character(haart2[, 'init.reg'])
row.reg2<-strsplit(haart2[, 'init.reg'], ",")
user.reg2<-sapply(all.reg,function(j) sapply(row.reg2,function(i) j %in% i))
haart2<-cbind(haart2,user.reg2)
haart<-rbind(haart,haart2)
head(haart,n=5)
```

```
##   male age aids cd4baseline logvl weight hemoglobin init.reg
## 1    1  25   0         NA    NA      NA          NA 3TC,AZT,EFV
## 2    1  49   0        143    NA 58.0608         11 3TC,AZT,EFV
## 3    1  42   1        102    NA 48.0816          1 3TC,AZT,EFV
## 4    0  33   0        107    NA 46.0000         NA 3TC,AZT,NVP
## 5    1  27   0         52     4     NA          NA 3TC,D4T,EFV
##   init.date last.visit death date.death death.within.1year fut loss.tfu
## 1 2003-07-01 2007-02-26     0      <NA>                 NA 365      1
## 2 2004-11-23 2008-02-22     0      <NA>                 NA 365      1
## 3 2003-04-30 2005-11-21     1 2006-01-11                 0 365      0
## 4 2006-03-25 2006-05-05     1 2006-05-07                 1  41      0
## 5 2004-09-01 2007-11-13     0      <NA>                 NA 365      1
##   3TC AZT EFV NVP D4T ABC DDI IDV LPV RTV SQV FTC TDF DDC NFV T20 ATV FPV
## 1    1    1    1    0    0    0    0    0    0    0    0    0    0    0    0    0    0
## 2    1    1    1    0    0    0    0    0    0    0    0    0    0    0    0    0    0
## 3    1    1    1    0    0    0    0    0    0    0    0    0    0    0    0    0    0
## 4    1    1    0    1    0    0    0    0    0    0    0    0    0    0    0    0    0
## 5    1    0    1    0    1    0    0    0    0    0    0    0    0    0    0    0    0
```

```
tail(haart,n=5)
```

```
##   male age aids cd4baseline logvl weight hemoglobin
## 1000  0 40.00000    1        131    NA 46.2672          8
## 1001  0 27.00000    0        232    NA      NA          NA
## 1002  1 38.72142    0        170    NA 84.0000          NA
## 1003  1 23.00000   NA        154 3.995635 65.5000         14
## 1004  0 31.00000    0        236    NA 45.8136          NA
##   init.reg init.date last.visit death date.death death.within.1year
## 1000 3TC,D4T,NVP 2003-07-03 2008-02-29     0      <NA>                 NA
## 1001 3TC,AZT,NVP 2003-12-01 2004-01-05     0      <NA>                 NA
## 1002 3TC,AZT,NVP 2002-09-26 2004-03-29     0      <NA>                 NA
## 1003 3TC,DDI,EFV 2007-01-31 2007-04-16     0      <NA>                 NA
## 1004 3TC,D4T,NVP 2003-12-03 2007-10-11     0      <NA>                 NA
##   fut loss.tfu 3TC AZT EFV NVP D4T ABC DDI IDV LPV RTV SQV FTC TDF DDC
## 1000 365      1    1    0    0    1    1    0    0    0    0    0    0    0    0
## 1001  35      0    1    1    0    1    0    0    0    0    0    0    0    0    0
## 1002 365      1    1    1    0    1    0    0    0    0    0    0    0    0    0
```

```
## 1003 75      0 1 0 1 0 0 0 1 0 0 0 0 0 0 0
## 1004 365     1 1 0 0 1 1 0 0 0 0 0 0 0 0 0
##      NFV T20 ATV FPV
## 1000  0  0  0  0
## 1001  0  0  0  0
## 1002  0  0  0  0
## 1003  0  0  0  0
## 1004  0  0  0  0
```

Question 2

14 points

Use the following code to generate data for patients with repeated measures of A1C (a test for levels of blood glucose).

```
genData <- function(n) {
  if(exists(".Random.seed", envir = .GlobalEnv)) {
    save.seed <- get(".Random.seed", envir = .GlobalEnv)
    on.exit(assign(".Random.seed", save.seed, envir = .GlobalEnv))
  } else {
    on.exit(rm(".Random.seed", envir = .GlobalEnv))
  }
  set.seed(n)
  subj <- ceiling(n / 10)
  id <- sample(subj, n, replace=TRUE)
  times <- as.integer(difftime(as.POSIXct("2005-01-01"), as.POSIXct("2000-01-01"), units='secs'))
  dt <- as.POSIXct(sample(times, n), origin='2000-01-01')
  mu <- runif(subj, 4, 10)
  a1c <- unsplit(mapply(rnorm, tabulate(id), mu, SIMPLIFY=FALSE), id)
  data.frame(id, dt, a1c)
}
x <- genData(500)
```

Perform the following manipulations: (2 points each)

1. Order the data set by id and dt.

```
x<-x[order(x[, 'id'], x[, 'dt']),]
```

2. For each id, determine if there is more than a one year gap in between observations. Add a new row at the one year mark, with the a1c value set to missing. A two year gap would require two new rows, and so forth.

```
library(lubridate)
```

```
##
## Attaching package: 'lubridate'
## The following object is masked from 'package:base':
##
##      date
d<-0
for (i in 1:499) {
  if (x[, 'id'][i]==x[, 'id'][i+1]) {
    d[i]<-floor(difftime(x[, 'dt'][i+1], x[, 'dt'][i])/365)
```

```

}
}
a<-data.frame(x$id[which(d==1)],x$dt[which(d==1)]+years(1),NA)
names(a)<-c('id','dt','a1c')
b<-data.frame(x$id[which(d==2)],x$dt[which(d==2)]+years(1),NA)
names(b)<-c('id','dt','a1c')
c<-data.frame(x$id[which(d==2)],x$dt[which(d==2)]+years(2),NA)
names(c)<-c('id','dt','a1c')
x<-rbind(x,a,b,c)
x<-x[order(x$id,x$dt),]

```

3. Create a new column visit. For each id, add the visit number. This should be 1 to n where n is the number of observations for an individual. This should include the observations created with missing a1c values.

```

visit<-1
for (i in 1:555) {
  if (x$id[i]==x$id[i+1]) {
    visit[i+1]<-visit[i]+1
  } else {
    visit[i+1]<-1
  }
}
x[, 'visit']<-visit

```

4. For each id, replace missing values with the mean a1c value for that individual.

```

for (i in 1:nrow(x)) {
  if(is.na(x$a1c[i])) {
    x$a1c[i] <- mean(x$a1c[which(x$id == x$id[i])], na.rm = TRUE)
  }
}

```

5. Print mean a1c for each id.

```

attach(x)

## The following object is masked _by_ .GlobalEnv:
##
##      visit
m<-matrix(0,nrow=2,ncol=50)
for (i in 1:50) {
  m[1,i]<-unique(id)[i]
  m[2,i]<-mean(x[, 'a1c'] [id==i])
}
row.names(m)=c('id','mean a1c')
print(m)

##           [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
## id      1.000000 2.000000 3.000000 4.000000 5.000000 6.000000 7.000000
## mean a1c 4.063372 7.544643 6.75764 3.892127 9.512311 7.555965 9.161686
##           [,8]      [,9]     [,10]     [,11]     [,12]     [,13]
## id      8.000000 9.000000 10.000000 11.000000 12.000000 13.000000
## mean a1c 7.189064 9.283873 7.975217 6.917562 7.034021 9.145282
##           [,14]     [,15]     [,16]     [,17]     [,18]     [,19]
## id     14.000000 15.000000 16.000000 17.000000 18.000000 19.000000

```

```
## mean a1c  6.623756  8.012406  4.222158  3.996034  9.164873  5.50721
##          [,20]    [,21]    [,22]    [,23]    [,24]    [,25]
## id       20.000000 21.000000 22.000000 23.000000 24.000000 25.000000
## mean a1c  3.726675  8.140939  5.637501  7.366889  7.439316  6.877135
##          [,26]    [,27]    [,28]    [,29]    [,30]    [,31]
## id       26.000000 27.000000 28.000000 29.000000 30.000000 31.000000
## mean a1c  6.556759  4.926457  7.433917  4.508086  6.045577  7.116586
##          [,32]    [,33]    [,34]    [,35]    [,36]    [,37]
## id       32.000000 33.000000 34.000000 35.0000  36.00000  37.000000
## mean a1c  6.568791  6.494069  6.768615  8.4767   9.60441  9.606253
##          [,38]    [,39]    [,40]    [,41]    [,42]    [,43]
## id       38.000000 39.000000 40.000000 41.000000 42.00000  43.000000
## mean a1c  5.355979  6.917013  9.530136  9.802424  3.89177  6.095849
##          [,44]    [,45]    [,46]    [,47]    [,48]    [,49]
## id       44.00000  45.000000 46.000000 47.000000 48.0000  49.000000
## mean a1c  9.09167  6.737204  9.621763  9.231489  6.4046  6.096076
##          [,50]
## id       50.000000
## mean a1c  8.962319
```

6. Print total number of visits for each id.

```
attach(x)
```

```
## The following object is masked _by_ .GlobalEnv:
##
##      visit
##
## The following objects are masked from x (pos = 3):
##
##      a1c, dt, id, visit
```

```
v<-matrix(0,nrow=2,ncol=50)
for (i in 1:50) {
  v[1,i]<-unique(id)[i]
  v[2,i]<-max(visit[id==i])
}
row.names(v)=c('id','max visit')
print(v)
```

```
##          [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12]
## id       1    2    3    4    5    6    7    8    9   10   11   12
## max visit 11   20   14   12   14   10    9   12   11   12   10   10
##          [,13] [,14] [,15] [,16] [,17] [,18] [,19] [,20] [,21] [,22]
## id       13   14   15   16   17   18   19   20   21   22
## max visit  8   12    8    9   12   10   10    9   10    8
##          [,23] [,24] [,25] [,26] [,27] [,28] [,29] [,30] [,31] [,32]
## id       23   24   25   26   27   28   29   30   31   32
## max visit  8   15   12   14   11   14   10    7   11    5
##          [,33] [,34] [,35] [,36] [,37] [,38] [,39] [,40] [,41] [,42]
## id       33   34   35   36   37   38   39   40   41   42
## max visit  8   12   11    9   17   15    8    7   17   14
##          [,43] [,44] [,45] [,46] [,47] [,48] [,49] [,50]
## id       43   44   45   46   47   48   49   50
## max visit 11   11   14    9   12   11   12   10
```

7. Print the observations for id = 15.

```
print(x[x$id==15,])
```

```
##      id      dt      a1c visit
## 11   15 2000-04-30 00:34:50 7.527105    1
## 406  15 2001-01-17 21:11:02 5.898371    2
## 306  15 2001-04-25 06:23:05 8.566593    3
## 1101 15 2002-04-25 06:23:05 8.012406    4
## 1102 15 2003-04-25 06:23:05 8.012406    5
## 484  15 2003-06-06 14:06:00 9.133769    6
## 1810 15 2004-06-06 14:06:00 8.012406    7
## 263  15 2004-08-20 17:47:11 8.936190    8
```

Question 3

10 points

Import the `addr.txt` file from the GitHub repository. This file contains a listing of names and addresses (thanks google). Parse each line to create a data.frame with the following columns: lastname, firstname, streetno, streetname, city, state, zip. Keep middle initials or abbreviated names in the firstname column. Print out the entire data.frame.

```
addr<-read.table('addr.txt',header=F,sep='\t',colClasses=c('character'))
u<-unlist(strsplit(addr[,1],split=" "))
u<-gsub("^\\s+|\\s+$","",u)
u<-u[u!=""]
x<-matrix(u,ncol=6,byrow=T)
y<-data.frame(streetno<-sub("^((\\w+)\\s?(.*)$)","\\1",x[,3]),streetname<-sub("^((\\w+)\\s?(.*)$)","\\2",x[,4]),
city<-sub("^((\\w+)\\s?(.*)$)","\\3",x[,5]),state<-sub("^((\\w+)\\s?(.*)$)","\\4",x[,6]),zip<-sub("^((\\w+)\\s?(.*)$)","\\5",x[,7]))
z<-as.data.frame(cbind(y,x)[-5])
colnames(z)<-c('streetno','streetname','lastname','firstname','city','state','zip')
z<-z[,c(3,4,1,2,5,6,7)]
print(z)
```

```
##      lastname  firstname streetno      streetname      city state
## 1      Bania   Thomas M.    725   Commonwealth Ave.    Boston   MA
## 2     Barnaby    David     373      W. Geneva St.    Wms. Bay   WI
## 3      Bausch    Judy     373      W. Geneva St.    Wms. Bay   WI
## 4     Bolatto   Alberto    725   Commonwealth Ave.    Boston   MA
## 5   Carlstrom    John     933      E. 56th St.    Chicago   IL
## 6 Chamberlin  Richard A.    111     Nowelo St.      Hilo     HI
## 7       Chuss    Dave    2145   Sheridan Rd     Evanston  IL
## 8       Davis    E. J.     933      E. 56th St.    Chicago   IL
## 9      Depoy    Darren    174      W. 18th Ave.    Columbus  OH
## 10      Griffin  Greg    5000   Forbes Ave.    Pittsburgh PA
## 11   Halvorsen   Nils     933      E. 56th St.    Chicago   IL
## 12      Harper    Al     373      W. Geneva St.    Wms. Bay   WI
## 13      Huang   Maohai    725 W. Commonwealth Ave.    Boston   MA
## 14     Ingalls   James G.    725 W. Commonwealth Ave.    Boston   MA
## 15     Jackson   James M.    725 W. Commonwealth Ave.    Boston   MA
## 16     Knudsen   Scott     373      W. Geneva St.    Wms. Bay   WI
## 17      Kovac    John    5640      S. Ellis Ave.    Chicago   IL
## 18   Landsberg   Randy    5640      S. Ellis Ave.    Chicago   IL
## 19          Lo   Kwok-Yung  1002      W. Green St.    Urbana    IL
## 20 Loewenstein  Robert F.    373      W. Geneva St.    Wms. Bay   WI
## 21      Lynch    John    4201     Wilson Blvd   Arlington  VA
```


## 22	Martini	Paul	174	W. 18th Ave.	Columbus	OH
## 23	Meyer	Stephan	933	E. 56th St.	Chicago	IL
## 24	Mrozek	Fred	373	W. Geneva St.	Wms. Bay	WI
## 25	Newcomb	Matt	5000	Forbes Ave.	Pittsburgh	PA
## 26	Novak	Giles	2145	Sheridan Rd	Evanston	IL
## 27	Odalen	Nancy	373	W. Geneva St.	Wms. Bay	WI
## 28	Pernic	Dave	373	W. Geneva St.	Wms. Bay	WI
## 29	Pernic	Bob	373	W. Geneva St.	Wms. Bay	WI
## 30	Peterson	Jeffrey	5000	Forbes Ave.	Pittsburgh	PA
## 31	Pryke	Clem	933	E. 56th St.	Chicago	IL
## 32	Rebull	Luisa	5640	S. Ellis Ave.	Chicago	IL
## 33	Renbarger	Thomas	2145	Sheridan Rd	Evanston	IL
## 34	Rottman	Joe	8730	W. Mountain View Ln	Littleton	CO
## 35	Schartman	Ethan	933	E. 56th St.	Chicago	IL
## 36	Spotz	Bob	373	W. Geneva St.	Wms. Bay	WI
## 37	Thoma	Mark	373	W. Geneva St.	Wms. Bay	WI
## 38	Walker	Chris	933	N. Cherry St.	Tucson	AZ
## 39	Wehrer	Cheryl	5000	Forbes Ave.	Pittsburgh	PA
## 40	Wirth	Jesse	373	W. Geneva St.	Wms. Bay	WI
## 41	Wright	Greg	791	Holmdel-Keyport Rd.	Holmdel	NY
## 42	Zingale	Michael	5640	S. Ellis Ave.	Chicago	IL
##	zip					
## 1	02215					
## 2	53191					
## 3	53191					
## 4	02215					
## 5	60637					
## 6	96720					
## 7	60208-3112					
## 8	60637					
## 9	43210					
## 10	15213					
## 11	60637					
## 12	53191					
## 13	02215					
## 14	02215					
## 15	02215					
## 16	53191					
## 17	60637					
## 18	60637					
## 19	61801					
## 20	53191					
## 21	22230					
## 22	43210					
## 23	60637					
## 24	53191					
## 25	15213					
## 26	60208-3112					
## 27	53191					
## 28	53191					
## 29	53191					
## 30	15213					
## 31	60637					
## 32	60637					

```
## 33 60208-3112
## 34      80125
## 35      60637
## 36      53191
## 37      53191
## 38      85721
## 39      15213
## 40      53191
## 41 07733-1988
## 42      60637
```

Question 4

2 points

The first argument to most functions that fit linear models are formulas. The following example defines the response variable `death` and allows the model to incorporate all other variables as terms. `.` is used to mean all columns not otherwise in the formula.

```
url <- "https://github.com/fonnesbeck/Bios6301/raw/master/datasets/haart.csv"
haart_df <- read.csv(url)[,c('death', 'weight', 'hemoglobin', 'cd4baseline')]
coef(summary(glm(death ~ ., data=haart_df, family=binomial(logit))))
```

```
##              Estimate Std. Error  z value    Pr(>|z|)
## (Intercept)  3.576411744 1.226870535  2.915069 0.0035561039
## weight      -0.046210552 0.022556001 -2.048703 0.0404911395
## hemoglobin   -0.350642786 0.105064078 -3.337418 0.0008456055
## cd4baseline  0.002092582 0.001811959  1.154872 0.2481427160
```

Now imagine running the above several times, but with a different response and data set each time. Here's a function:

```
myfun <- function(dat, response) {
  form <- as.formula(response ~ .)
  coef(summary(glm(form, data=dat, family=binomial(logit))))
}
```

Unfortunately, it doesn't work. `tryCatch` is "catching" the error so that this file can be knit to PDF.

```
tryCatch(myfun(haart_df, death), error = function(e) e)
```

```
## <simpleError in model.frame.default(formula = form, data = dat, drop.unused.levels = TRUE): variable
```

What do you think is going on? Consider using `debug` to trace the problem.

The error is that `death` itself is not a character object variable that needs to be in the `as.formula` function. To fix it, we need to paste it with `"~ ."`.

5 bonus points

Create a working function.

```
myfun <- function(dat, response) {
  form <- as.formula(paste(response, "~ ."))
  coef(summary(glm(form, data=dat, family=binomial(logit))))
}
myfun(haart_df, "death")
```

```
##              Estimate Std. Error  z value    Pr(>|z|)
```

```
## (Intercept)  3.576411744  1.226870535  2.915069  0.0035561039
## weight      -0.046210552  0.022556001 -2.048703  0.0404911395
## hemoglobin  -0.350642786  0.105064078 -3.337418  0.0008456055
## cd4baseline  0.002092582  0.001811959  1.154872  0.2481427160
```

This new function with paste, deparse and substitute to process the formula produces exactly the same results as before.

JC Grading +5